

Deep Learning

Practical 3

1. From last Practicals

The base toy model we used from last two practicals is:

```

import tensorflow as tf
import numpy as np

mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train, x_test = x_train / 255.0 , x_test / 255.0

input_layer = tf.keras.layers.Input(name="input_layer", shape=(28, 28))
flatten_layer = tf.keras.layers.Flatten()(input_layer)
hidden_layer = tf.keras.layers.Dense(units=128)(flatten_layer)
hidden_layer = tf.keras.layers.Activation("relu")(hidden_layer)
hidden_layer = tf.keras.layers.Dropout(0.2)(hidden_layer)
hidden_layer = tf.keras.layers.Dense(units=10)(hidden_layer)
output_layer = tf.keras.layers.Activation(
    "softmax", name="output_layer"
)(hidden_layer)

model = tf.keras.models.Model(inputs=[input_layer], outputs=[output_layer])

model.summary(150)

model.compile(
    optimizer = "Adam",
    loss = {
        "output_layer": tf.keras.losses.SparseCategoricalCrossentropy()
    },
    metrics = ["acc"],
)

model.fit(
    x_train,
    y_train,
    validation_data=(x_test, y_test),
    batch_size = 32,
    epochs = 5
)

```

The objectif of this practical is to transform it into an encoder/
Decoder Sequentially Separated Convolutional Residual Variadic Neural Network !

Let's proceed step by step.

Provide a working .py file for each part.

2. Convolutional Transformation

2.1. The Idea

The idea here is to transform the existing `dense` layers into `convolutional` layers. `Convolutional` layers are more suited for image data as they can extract spatial hierarchies of features from the images.

2.2. Methodology

1. Input Shape: Change the input shape to include the channel dimension. For MNIST, this would be (28, 28, 1) as the images are grayscale.
2. `Convolutional` Layers: Replace `dense` layers with `convolutional` layers (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D). You'll still need a `Flatten` layer before the final `Dense` layer to convert the multi-dimensional `convolutional` output to a flat form.

Replace the first `dense` layer with a `convolutional` layer with `kernel_size = (1, 1)` and `filters = 128`. For now, keep the `Flatten` and final `Dense` layer.

Add 3 more `convolutions`, Use 7×7 kernels, 128 filters and `padding = "same"` for this layers.

3. Variadic Transformation

3.1. The Idea

The goal is to make the network capable of handling images of `varying` sizes.

3.2. Methodology

1. Input Shape: Change the input shape to have `None` for the width and height dimensions, allowing for `variability` in image size.
2. `GlobalAveragePooling2D` (https://www.tensorflow.org/api_docs/python/tf/keras/layers/GlobalAveragePooling2D): This layer will help to transform `variably` shaped inputs into a fixed-size vector, making it possible to output a fixed-size probability vector.

Change input shape to `(None, None, 1)`.

Change the final `Dense` layer by a `Conv2D` layers with 10 filters. Include a `GlobalAveragePooling2D` layer just after to have a 10-sized vector.

4. Residual

4.1. The Idea

The concept behind adding `residual` layers is to enable the learning of identity functions, which aids in addressing the vanishing gradient problem in deep neural networks. By incorporating `residuals`, or “shortcut connections”, around each `convolutional` block, the network becomes fully `residual`. These connections allow the network to preserve information from earlier layers and pass it through deeper layers without significant degradation.

4.2. Methodology

For each `convolutional` block, you'll introduce a shortcut that bypasses the block and adds its input directly to its output. This can be done as follows:

1. Before a `convolutional` block, save the input to that block.
2. After the `convolutional` block, add the saved input to the output of the block.

Here's a conceptual example:

```
input_to_block = x # Save input before the block
# Convolutional block
x = tf.keras.layers.Conv2D(...)(x)
x = tf.keras.layers.BatchNormalization()(x)
x = tf.keras.layers.Activation("relu")(x)

# Adding the residual
x = x + input_to_block
```

Repeat this pattern around each `convolutional` block to ensure the entire network is fully `residual` from top to bottom.

5. Sequentially Separated

5.1. The Idea

Sequentially separated `convolutions` involve breaking down larger `convolutions` into smaller, more manageable parts, which can increase depth and reduce parameters.

5.2. Disclaimer

We will only be dealing with square odd-sized kernels, but the methodology is the same for any sized kernel, just do it separately for the two dimensions and use a 2-sized kernel for the final one in case of even-sized requests.

5.3. Methodology

1. Get the `receptive field` (φ) size you want to achieve (for example 9×9 gives $\varphi = 9$)
2. Compute the number of 3×3 `convolutions` needed to achieve such `receptive field`, using the formula: $n = \frac{\varphi-1}{2}$ (here $n = 4$)
3. Replace the φ -sized kernel `convolution` by n 3×3 -sized kernel `convolutions`. In our study case, replace

```
x = tf.keras.layers.Conv2D(filters=c, kernel_size=(9, 9), padding="same")(x)
```

by

```
for _ in range((9-1)//2):
    x = tf.keras.layers.Conv2D(filters=c, kernel_size=(3, 3), padding="same")(x)
```

4. Add activations, normalizations and regularization as you see fit. For example:

```
for _ in range((9-1)//2):
    x = tf.keras.layers.Conv2D(filters=c, kernel_size=(3, 3), padding="same")(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.Activation("tanh")(x)
    x = tf.keras.layers.Dropout(0.2)(x)
```

5. Replace all 3×3 -sized `convolution` by a 3×1 followed by a 1×3 :

```
for _ in range((9-1)//2):
    x = tf.keras.layers.Conv2D(filters=c, kernel_size=(3, 1), padding="same")(x)
    x = tf.keras.layers.Conv2D(filters=c, kernel_size=(1, 3), padding="same")(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.Activation("tanh")(x)
    x = tf.keras.layers.Dropout(0.2)(x)
```

In your code, change all `convolutions` to Sequentialy Separated `convolutions` !

6. Encoder/Decoder Convolutions Blocks

6.1. The Idea

Apply a channel reduction transformation before applying the `sequential separation` and transform back to the original size at the end to sum for the `residual` block, thus lowering the number of parameters of the network.

6.2. Methodology

1. Add a 1×1 -sized kernel **convolution** at the beginning of each **convolution block**, with a low filters value in order to drastically lower the computation size (like $n = 32$).
3. Change all filters values in the **convolutions** that will be compressed to the chosen n .
2. Add a 1×1 -sized kernel **convolution** at the end of each **convolution block**, with the original number of filters value in order to get back the original number of channels and be able to add them together for the **residual** block.

Apply this methodology for each **sequentialy separated convolution block** in your code, try different value for n and see what works best.

7. Cifar-10

Update your code to run on the Cifar-10 dataset (https://www.tensorflow.org/api_docs/python/tf/keras/datasets/cifar10/load_data)